

# Assignment 3: Federated Learning

## 1 Coding Task

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.datasets as datasets
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
from torch.multiprocessing import Pool
from torch.utils.data import RandomSampler
import torch.multiprocessing as mp

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Hyperparameters
learning_rate = 0.1
batch_size = 128
num_communications = 50
num_workers = 2 # Number of models to train in parallel
num_local_steps = 5

# Loss function and optimizer
criterion = nn.CrossEntropyLoss()

# MNIST dataset
train_dataset = datasets.MNIST(root='./data', train=True,
                                transform=transforms.ToTensor(), download=True)
test_dataset = datasets.MNIST(root='./data', train=False,
                                transform=transforms.ToTensor())

# Neural network model
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(784, 512)
        self.fc2 = nn.Linear(512, 256)
        self.fc3 = nn.Linear(256, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = x.view(-1, 784)
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.fc3(x)

        return x

# Function to average the model weights
```

```

def average_models(models):
    model_params = [model.state_dict() for model in models]
    averaged_params = {}
    for param_name in model_params[0]:
        params = torch.stack([model_params[i][param_name] for i in range(len(models))])
        averaged_params[param_name] = torch.mean(params, dim=0)
    return averaged_params

# Function to duplicate a model
def duplicate_model(model, num_duplicates):
    model_dicts = [model.state_dict() for _ in range(num_duplicates)]
    duplicated_models = [Net().to(device) for _ in range(num_duplicates)]
    for i, model_dict in enumerate(model_dicts):
        duplicated_models[i].load_state_dict(model_dict)
    return duplicated_models

# Function to train a model
def train_model(args):
    model, optimizer, criterion, random_sampler, train_dataset = args
    model.to(device)

    # Create a new data loader with the random sampler
    train_loader_random = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=batch_size,
                                                       sampler=random_sampler)

    total_loss = 0
    model.train()
    for i, (images, labels) in enumerate(train_loader_random):
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()

    model.cpu()
    return model, total_loss

# Random sampler for train loader
train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
random_sampler = RandomSampler(train_dataset, replacement=True,
                               num_samples=num_local_steps * len(train_loader))

# Training Loop
losses_array = []
model = Net().to(device)
if __name__ == '__main__':
    mp.set_start_method('spawn')
    for _ in range(num_communications):
        models = duplicate_model(model, num_workers)

```

```

optimizers = [optim.SGD(model.parameters(), lr=learning_rate) for model in models]
args_list = [(models[i], optimizers[i], criterion, random_sampler, train_dataset) for i in range(num_workers)]

ctx = mp.get_context("spawn")
with ctx.Pool(processes=num_workers) as pool:
    results = pool.map(train_model, args_list)
models = [result[0].to(device) for result in results]
total_losses = [result[1] for result in results]
total_loss = sum(total_losses)
# Average the models' weights to create the ensemble model
ensemble_model_params = average_models(models)
model.load_state_dict(ensemble_model_params)
losses_array.append(total_loss)

plt.plot(losses_array)
plt.xlabel('Number of Communications')
plt.ylabel('Total Loss')
plt.title('Training Loss over Communications')
plt.savefig('CodingTask.png')
plt.show()

# Testing the Model
test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
                                          batch_size=batch_size, shuffle=False)

model.eval()
with torch.no_grad():
    correct = 0
    total = 0
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    print(f'Test Accuracy: {100 * correct / total}%',)

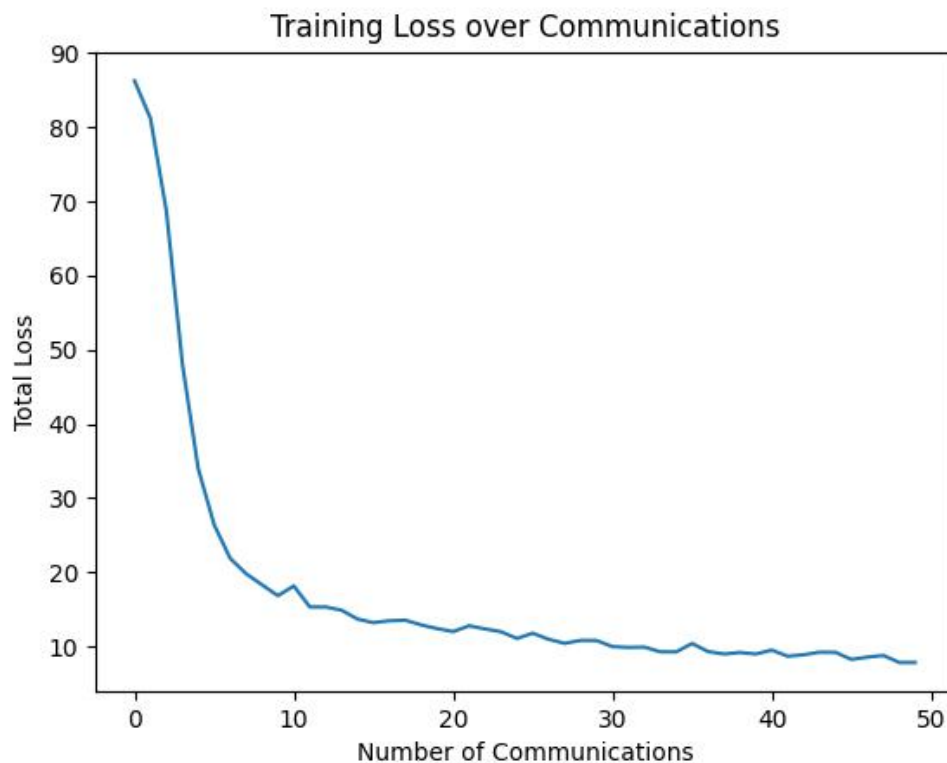
```

The screenshots of the output results are as follows.

```

(.conda) lumos@lumosl:~/Desktop/python/program/week2/assignment$ python CodingTask.py
Test Accuracy: 93.36%
(.conda) lumos@lumosl:~/Desktop/python/program/week2/assignment$ █

```



The original Coding Task was not designed to use the GPU. However, in order to accelerate the training process and ensure that each test set runs within approximately five minutes in my computer, the GPU was utilized. By leveraging the parallel processing capabilities of the GPU, the computational time for training the neural network models was significantly reduced.

## 2 Questions

In addition to the basic implementation of Federated Learning in the Coding Task, four new models with specified hyperparameter configurations were trained using the same code framework.

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.datasets as datasets
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
from torch.utils.data import DataLoader, RandomSampler
import torch.multiprocessing as mp

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
# Define a simple neural network model for MNIST classification
```

```
class Net(nn.Module):
```

```
    def __init__(self):
```

```
        super(Net, self).__init__()
```

```
        self.fc1 = nn.Linear(784, 512)
```

```
        self.fc2 = nn.Linear(512, 256)
```

```
        self.fc3 = nn.Linear(256, 10)
```

```
        self.relu = nn.ReLU()
```

```
    def forward(self, x):
```

```
        x = x.view(-1, 784)
```

```
        x = self.relu(self.fc1(x))
```

```
        x = self.relu(self.fc2(x))
```

```
        x = self.fc3(x)
```

```
        return x
```

```
def average_models(models):
```

```
    model_params = [model.state_dict() for model in models]
```

```
    averaged_params = {}
```

```
    for param_name in model_params[0]:
```

```
        params = torch.stack([model_params[i][param_name] for i in range(len(models))])
```

```
        averaged_params[param_name] = torch.mean(params, dim=0)
```

```
    return averaged_params
```

```
def duplicate_model(model, num_duplicates):
```

```
    model_dicts = [model.state_dict() for _ in range(num_duplicates)]
```

```
    duplicated_models = [Net().to(device) for _ in range(num_duplicates)]
```

```
    for i, model_dict in enumerate(model_dicts):
```

```
        duplicated_models[i].load_state_dict(model_dict)
```

```
    return duplicated_models
```

```
# Function to train a single model on the Local dataset
```

```
def train_model(args):
```

```
    model, optimizer, criterion, random_sampler, train_dataset, batch_size = args
```

```
    model.to(device) # Ensure the model is on the correct device (GPU if available)
```

```
    train_loader_random = DataLoader(dataset=train_dataset, batch_size=batch_size, sampler=random_sampler)
```

```
    total_loss = 0
```

```
    model.train()
```

```
    for i, (images, labels) in enumerate(train_loader_random):
```

```
        images, labels = images.to(device), labels.to(device)
```

```
        optimizer.zero_grad()
```

```
        outputs = model(images)
```

```
        loss = criterion(outputs, labels)
```

```

        loss.backward()
        optimizer.step()
        total_loss += loss.item()
    model.cpu()
    return model, total_loss

def test_model(model, test_loader):
    model.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            model.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    model.cpu()
    return correct / total

def federated_training(batch_size, num_communications, learning_rate, num_workers, num_local_steps):
    train_dataset = datasets.MNIST(root='./data', train=True, transform=transforms.ToTensor(), download=True)
    test_dataset = datasets.MNIST(root='./data', train=False, transform=transforms.ToTensor())
    train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
    test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)
    random_sampler = RandomSampler(train_dataset, replacement=True, num_samples=num_local_steps * len(train_loader))
    criterion = nn.CrossEntropyLoss()
    model = Net().to(device) # Ensure the model starts on the correct device
    losses = []
    train_accuracies = []
    test_accuracies = []
    for _ in range(num_communications):
        models = duplicate_model(model, num_workers)
        optimizers = [optim.SGD(model.parameters(), lr=learning_rate) for model in models]
        args_list = [(models[i], optimizers[i], criterion, random_sampler, train_dataset, batch_size)

for i in
                    range(num_workers)]
    # 使用 spawn 上下文创建 Pool
    ctx = mp.get_context("spawn")
    with ctx.Pool(processes=num_workers) as pool:

```

```

        results = pool.map(train_model, args_list)
        models = [result[0].to(device) for result in results]
        total_losses = [result[1] for result in results]
        total_loss = sum(total_losses)
        losses.append(total_loss)
        ensemble_model_params = average_models(models)
        model.load_state_dict(ensemble_model_params)
        train_accuracy = test_model(model, train_loader)
        test_accuracy = test_model(model, test_loader)
        train_accuracies.append(train_accuracy)
        test_accuracies.append(test_accuracy)
    return losses, train_accuracies, test_accuracies

if __name__ == '__main__':
    if mp.get_start_method(allow_none=True) is None:
        mp.set_start_method('spawn')
    hyperparam_configs = [
        (32, 50, 0.1, 2, 5),
        (32, 50, 0.1, 4, 5),
        (32, 50, 0.1, 8, 5),
        (32, 50, 0.1, 4, 20)
    ]
    for i, config in enumerate(hyperparam_configs):
        batch_size, num_communications, learning_rate, num_workers, num_local_steps = config
        print(f"Running experiment with config: batch size={batch_size}, num communications={num_communications}, "
              f"learning rate={learning_rate}, num workers={num_workers}, num local steps={num_local_steps}")
        losses, train_accuracies, test_accuracies = federated_training(batch_size, num_communications,
                                                                    learning_rate,
                                                                    num_workers, num_local_steps)

        plt.figure()
        plt.plot(range(num_communications), losses)
        plt.xlabel('Number of Communications')
        plt.ylabel('Training Loss')
        plt.title(f'Training Loss vs Communications (Config {i + 1})')
        filename = f'training_loss_config_{i + 1}.png'
        plt.savefig(filename)
        plt.show()
        final_train_accuracy = train_accuracies[-1]
        final_test_accuracy = test_accuracies[-1]
        print(f"Final Training Accuracy: {final_train_accuracy * 100:.2f}%")
        print(f"Final Test Accuracy: {final_test_accuracy * 100:.2f}%")
        print("-" * 50)

```


The screenshots of the output results are presented below.

```
问题 1 输出 调试控制台 终端 端口 JUPYTER 评论
(.conda) lumos@lumosl:~/Desktop/python/program/week2/assignment$ python main.py
Running experiment with config: batch size=32, num communications=50, learning rate=0.1, num workers=2, num local steps=5
Final Training Accuracy: 99.82%
Final Test Accuracy: 98.14%
-----
Running experiment with config: batch size=32, num communications=50, learning rate=0.1, num workers=4, num local steps=5
Final Training Accuracy: 99.93%
Final Test Accuracy: 98.20%
-----
Running experiment with config: batch size=32, num communications=50, learning rate=0.1, num workers=8, num local steps=5
Final Training Accuracy: 99.95%
Final Test Accuracy: 98.29%
-----
Running experiment with config: batch size=32, num communications=50, learning rate=0.1, num workers=4, num local steps=20
Final Training Accuracy: 100.00%
Final Test Accuracy: 98.37%
-----
(.conda) lumos@lumosl:~/Desktop/python/program/week2/assignment$
```

Here is the table of training and test accuracies of different hyperparameter configurations.


Batch Size	Num Communications	Learning Rate	Num Workers	Num Local Steps	Final Training Accuracy	Final Test Accuracy
32	50	0.1	2	5	99.82	98.14
32	50	0.1	4	5	99.93	98.2
32	50	0.1	8	5	99.95	98.29
32	50	0.1	4	20	100.0	98.37

The plots of different hyperparameter configurations are shown as follows.





Training Loss vs Communications (Config 2)



Training Loss vs Communications (Config 3)

